

Higgs Boson Machine Learning Challenge



DataUniversalis

Miaozhi YU, Shuo ZHANG, Bin FANG, Chuan SUN

Aug. 30, 2016

Outline

❑ **Workflow**

❑ Exploratory Data Analysis

❑ Models


❑ Random Forests

❑ Gradient Boosted Model

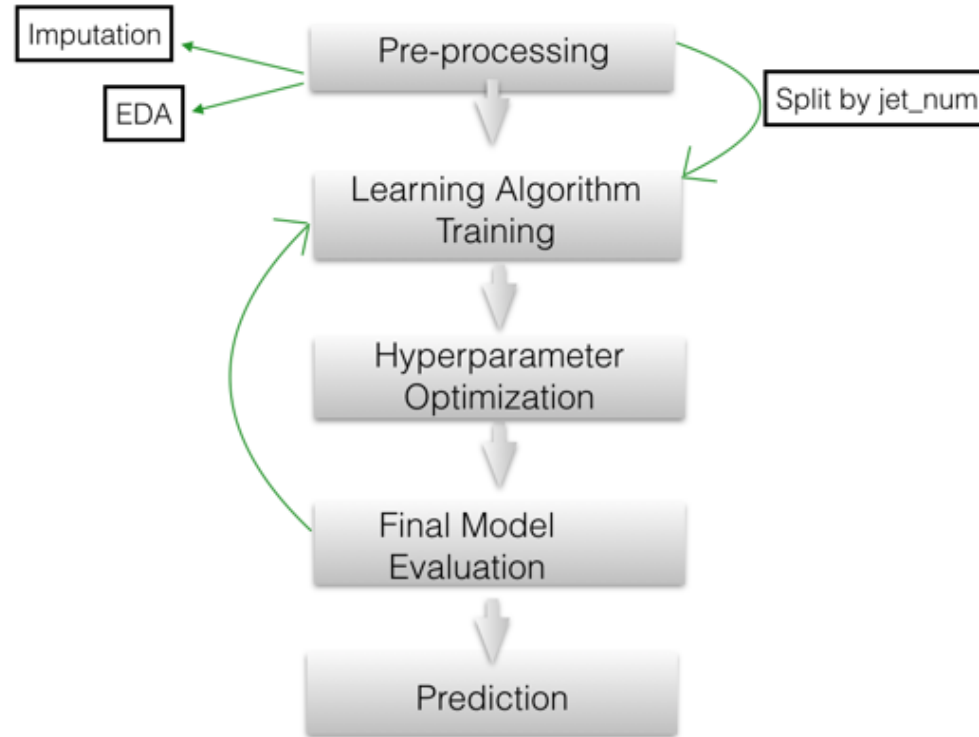
❑ Neural Networks

❑ XGBoost

❑ Lesson Learned




Workflow




EDA -- Missingness Imputation by kNN


PRI_jet_num = 0




PRI_jet_num = 1




PRI_jet_num = 2 or 3



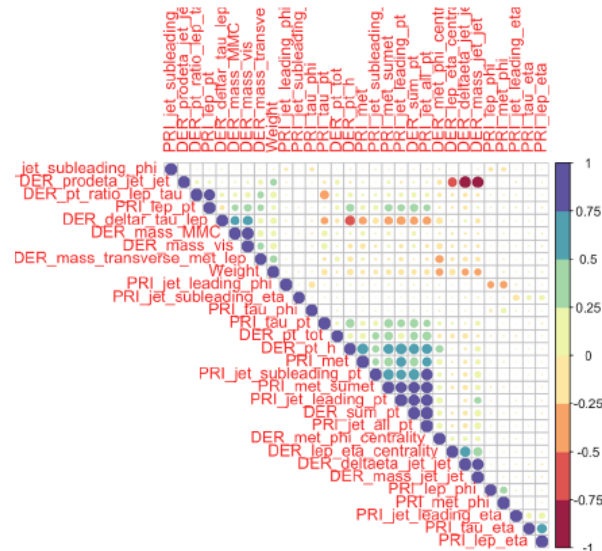
EDA -- Correlation



PRI_jet_num = 0



PRI_jet_num = 1



PRI_jet_num = 2 or 3

- ❑ Variables of same prefix (with some uppercase letters) have high positive or negative correlation.

EDA -- Principal Component Analysis


Principal Components Analysis

Call: principal(r = df0, nfactors = 7, rotate = "none")

Standardized loadings (pattern matrix) based upon correlation matrix

	PC1	PC2	PC3	PC4	PC5	PC6	PC7	h2	u2	com
DER_mass_MMC	0.90	-0.18	-0.11	0.21	-0.04	0.00	-0.01	0.90	0.103	1.2
DER_mass_transverse_met_lep	0.21	0.80	-0.46	-0.05	0.04	0.00	0.00	0.90	0.095	1.8
DER_mass_vis	0.88	-0.23	-0.12	0.12	-0.03	0.00	-0.01	0.85	0.145	1.2
DER_pt_h	0.18	0.44	0.82	0.16	-0.04	-0.01	0.00	0.93	0.070	1.7
DER_deltar_tau_lep	0.51	-0.52	-0.03	0.55	-0.13	-0.01	-0.01	0.84	0.159	3.1
DER_pt_tot	0.18	0.44	0.82	0.16	-0.04	-0.01	0.00	0.93	0.070	1.7
DER_sum_pt	0.89	0.02	-0.09	-0.33	0.07	0.00	0.01	0.91	0.094	1.3
DER_pt_ratio_lep_tau	0.09	0.59	-0.44	0.59	-0.11	0.00	0.01	0.91	0.088	3.0
DER_met_phi_centrality	0.10	-0.07	0.58	0.39	-0.10	-0.01	-0.01	0.52	0.484	1.9
PRI_tau_pt	0.64	-0.32	0.19	-0.62	0.12	0.00	0.01	0.95	0.047	2.7
PRI_tau_eta	0.01	-0.03	0.02	0.20	0.86	0.02	-0.01	0.77	0.225	1.1
PRI_tau_phi	-0.01	0.01	0.00	-0.01	-0.02	0.67	-0.60	0.81	0.190	2.0
PRI_lep_pt	0.76	0.37	-0.35	0.12	-0.02	0.00	0.01	0.85	0.151	2.0
PRI_lep_eta	0.01	-0.02	0.04	0.20	0.86	0.02	-0.01	0.77	0.225	1.1
PRI_lep_phi	0.00	-0.01	0.00	-0.01	0.02	-0.85	-0.03	0.72	0.278	1.0
PRI_met	0.12	0.77	0.02	-0.36	0.10	0.00	0.00	0.75	0.254	1.5
PRI_met_phi	0.00	-0.01	0.02	0.01	0.00	0.46	0.82	0.89	0.114	1.6
PRI_met_sumet	0.47	0.14	0.47	-0.14	0.02	0.01	-0.01	0.48	0.518	2.4

Parallel Analysis Scree Plots



❑ Importance: mass-related variables.

❑ Impute missing values:

DER_mass_MMC

Outline

- ❑ Workflow
- ❑ Exploratory Data Analysis
- ❑ **Models**
 - ❑ **Random Forests**
 - ❑ Gradient Boosted Model
 - ❑ Neural Networks
 - ❑ XGBoost
- ❑ Lesson Learned





Random Forests



- ❑ Two parameters: `ntree` and `mtry`.
- ❑ 1st tuning: `ntree = c(2000,5000,8000)` `mtry = c(3,4,5,6)`
- ❑ However, R crashed due to large computation
- ❑ 2nd tuning: `ntree = c(500,800,1000)` `mtry = c(3,4,5)`





Random Forests(DF:PRI_jet_num=0, ntree = 1000)





Random Forests(DF:PRI_jet_num=1, ntree = 1000)



Random Forests(DF:PRI_jet_num=2&3, ntree = 1000)



Random Forests(PRI_jet_num=2&3, ntree = 2 & 10)



Random Forests---Summary

Pros:

- ❑ Theoretically, it can never overfit
- ❑ Good with very large data set
- ❑ Robust against outliers

Cons:

- ❑ Not a fast algorithm. Very time and memory space consuming
- ❑ Less accurate than boosted tree models
- ❑ Cannot handle NAs







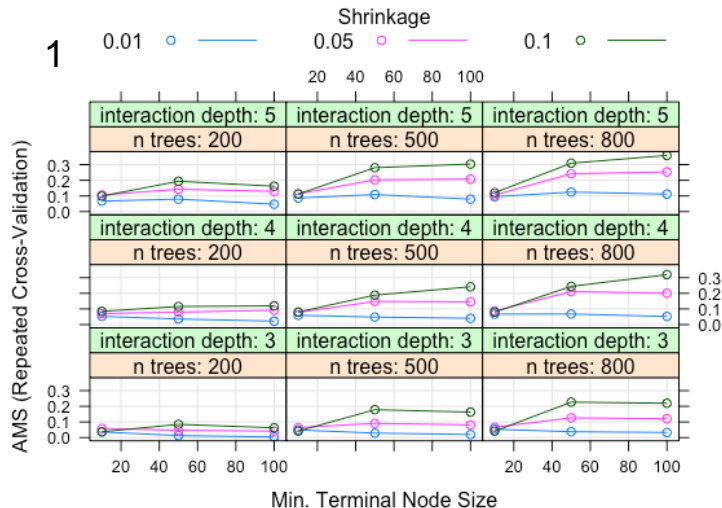
Outline

- ❑ Workflow
- ❑ Exploratory Data Analysis
- ❑ **Models**
 - ❑ Random Forests
 - ❑ **Gradient Boosted Model**
 - ❑ Neural Networks
 - ❑ XGBoost
- ❑ Lesson Learned



Gradient Boosting Model (DF:PRI_jet_num=1)

- ❑ Tuning four parameters, 5-fold validation
- ❑ 1st tuning: interaction.depth = c(3, 4, 5), n.trees = c(200, 500, 800), shrinkage = c(0.1, 0.05, 0.01), n.minobsinnode = c(10, 50, 100)
- ❑ 2nd tuning: 800 trees and interaction.depth (5,7,10)
- ❑ 3rd tuning: 5 interaction.depth and large trees(800, 2000, 5000)
- ❑ 4th tuning: 5 interaction.depth and larger trees(5000, 7500, 10000)




Tuning	AUC
1	0.88
2	0.86
3	0.91
4	0.93

Parameter tuning for the other 2 subsets

`jet_num=0`

`interaction.depth = 5, n.trees = 10000, shrinkage = 0.1,`
`n.minobsinnode = 100`


AUC=0.907



`jet_num=c(2,3)`

`interaction.depth = 5, n.trees = 800, shrinkage = 0.05,`
`n.minobsinnode = 100`

AUC=0.91



Summary of Gradient Boosting Model

Insight:

- ❑ Interaction.depth: **sqr(number of variables)**
- ❑ N.trees: **10000**, reduce overfitting by smaller learning rate and cross-validation
- ❑ Shrinkage: big trees: **0.001**; small trees: **0.1**

Pros:

- ❑ N.minobsinnode: **100**: reduce variance in predictions at leaves
- ❑ Use all features
- ❑ Higher accuracy and lower possibility of overfitting with more trees

Cons:

- ❑ Susceptible to outliers
- ❑ Lack of interpretability and higher complexity
- ❑ Harder to tune hyperparameters than other models
- ❑ Slow to train


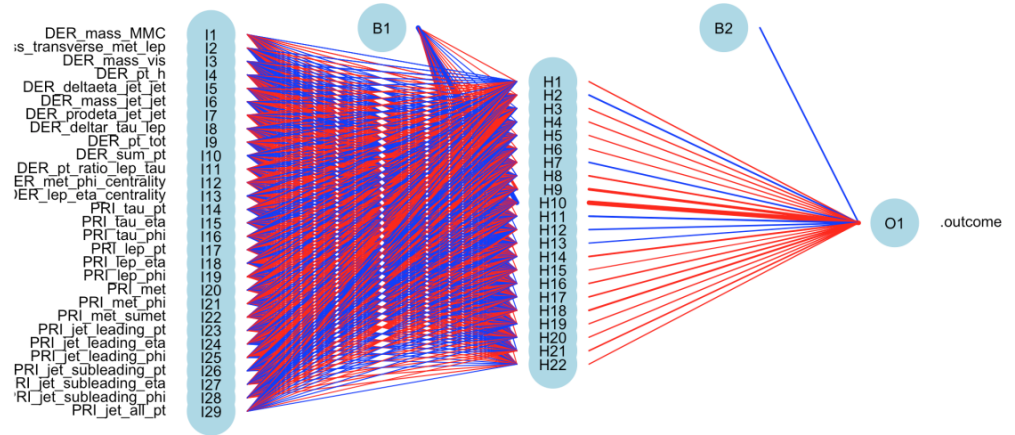
Outline

- ❑ Workflow
- ❑ Exploratory Data Analysis
- ❑ **Models**
 - ❑ Random Forests
 - ❑ Gradient Boosted Model
 - ❑ **Neural Networks**
 - ❑ XGBoost
- ❑ Lesson Learned

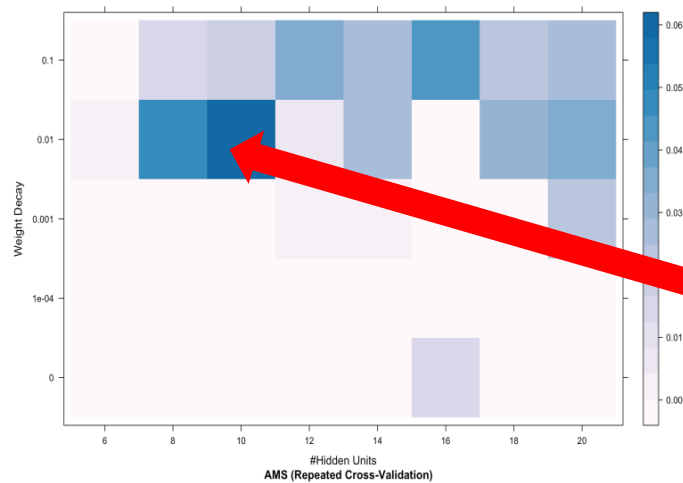


Neural Networks


- ❑ Single hidden layer NN model
“nnet” under “Caret” library.
- ❑ Two tuning parameters: number
of **hidden units** and **weight
decay**.
- ❑ Select initial combinations of
ranges of hidden units between
6-20, and decay weights
between **0-0.5**.




Neural Networks -- jet_number = 0



Best set:
hidden size = 10
weight decay = 0.01
Accuracy = 82.54%




hidden size = 10, weight decay = 0.01, accuracy = 82.54%





❑ Hidden units: seq(6,20,2); Weight decay: c(1e-04, 1e-03, 0.01, 0.1, 0).

❑ Fluctuation trend of AMS is noticed as hidden unit increases.

Neural Networks -- jet_number = 2/3



Best set:
hidden size = 22;
weight decay = 0.1
Accuracy = 80.4%



Hidden units: seq(6,22,2); Weight decay: c(1e-03,0.01,0.1,0.15).


AMS steadily increases as more hidden units used.

Neural Networks -- Summary

❑ Insight

- ❑ Best performance when weight decay = 0.1 or 0.15.
- ❑ More hidden nodes may be required.

❑ Pros & Cons

- ❑ Can handle hundreds or thousands of input variables.
 - ❑ Once trained, predictions are fast.
 - ❑ Training is computationally expensive.
 - ❑ Need a portion of the data for validation during NN training.
- 

Outline

- ❑ Workflow
- ❑ Exploratory Data Analysis
- ❑ **Models**
 - ❑ Random Forests
 - ❑ Gradient Boosted Model
 - ❑ Neural Networks
 - ❑ **XGBoost**
- ❑ Lesson Learned



XGboost

6 parameters to tune

nrounds

max_depth

eta

● Guidelines







- Begin with rough grid, gradually fine tuning
- Prefer slow learners (small eta, eg. 0.005)
- Tune max_depth (9 or 10)
- Fixed colsample_bytree (0.7)
- Fixed gamma (5)
- Fixed min_child_weight (0.1)

```
Parameter tuning for subset with jet number 0
(nrounds=c(20, 60, 150), max_depth=c(10, 15, 20), eta=c(0.001, 0.05), colsample_bytree=0.7, gamma=5, min_child_weight=0.1)
(nrounds=c(250), max_depth=c(7, 10), eta=c(0.001), colsample_bytree=0.7, gamma=c(1, 5), min_child_weight=0.1)
(nrounds=c(100, 150, 200), max_depth=c(8, 9, 10), eta=c(0.001, 0.005), colsample_bytree=0.7, gamma=c(5), min_child_weight=0.1)
(nrounds=c(60, 100, 150, 200, 400), max_depth=c(10), eta=c(0.001), colsample_bytree=0.7, gamma=c(5), min_child_weight=0.1)

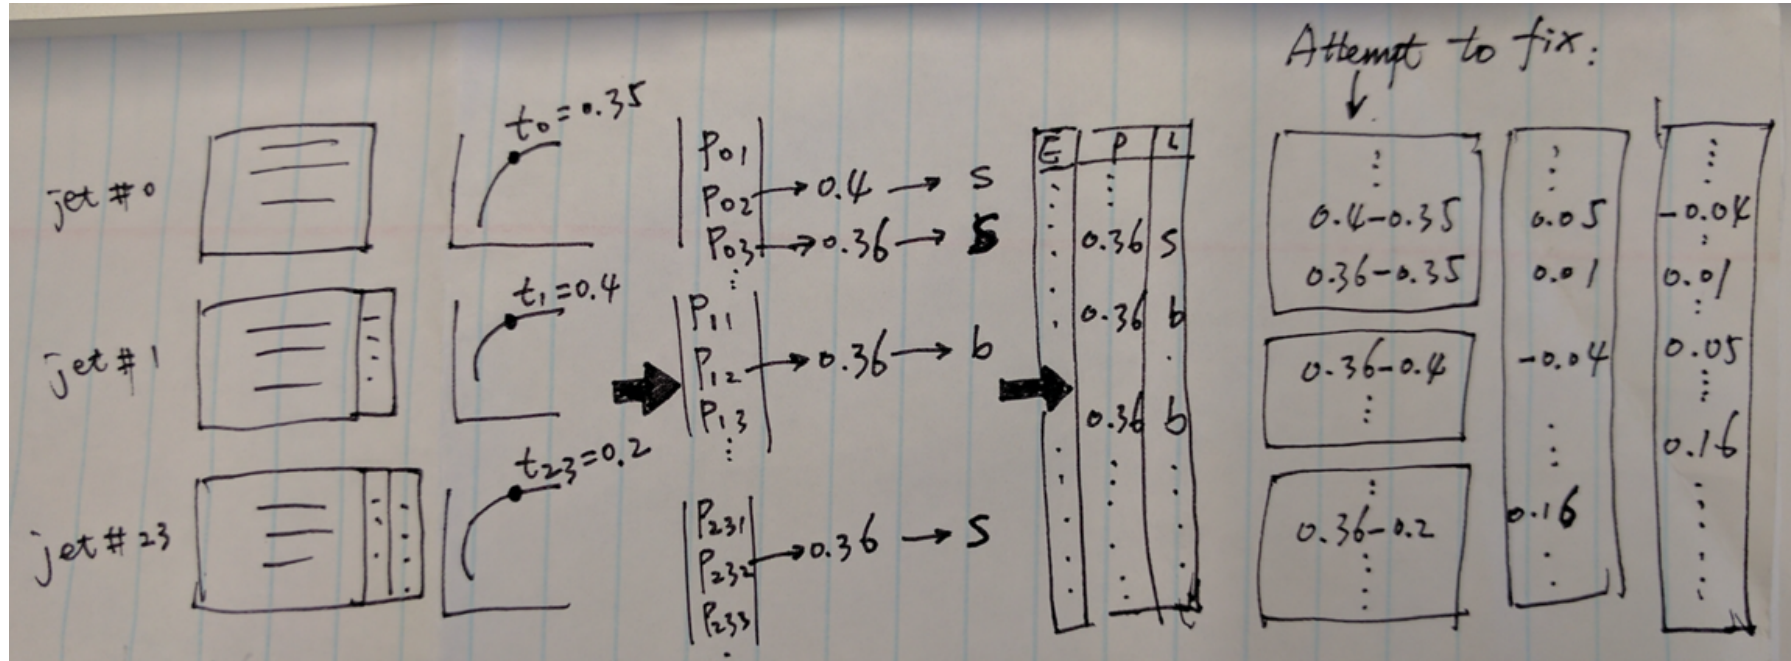
Parameter tuning for subset with jet number 1
(nrounds=c(20, 60, 150), max_depth=c(10, 15, 20), eta=c(0.001, 0.05), colsample_bytree=0.7, gamma=5, min_child_weight=0.1)
(nrounds=c(200, 300), max_depth=c(10), eta=c(0.05), colsample_bytree=0.7, gamma=5, min_child_weight=0.1)
(nrounds=c(200, 250), max_depth=c(8, 10, 12), eta=c(0.02, 0.05, 0.15), colsample_bytree=0.7, gamma=5, min_child_weight=0.1)

Parameter tuning for subset with jet number 2&3
(nrounds=c(20, 60, 150), max_depth=c(10, 15, 20), eta=c(0.001, 0.05), colsample_bytree=0.7, gamma=5, min_child_weight=0.1)
(nrounds=c(200, 400), max_depth=c(20), eta=c(0.001, 0.05), colsample_bytree=0.7, gamma=5, min_child_weight=0.1)
(nrounds=c(200, 250), max_depth=c(10, 20), eta=c(0.001, 0.003), colsample_bytree=0.7, gamma=5, min_child_weight=0.1)
```


Parameter tuning for all 3 subsets



When combining all predictions ...



Outline

- ❑ Workflow
- ❑ Exploratory Data Analysis
- ❑ Models
 - ❑ Random Forests
 - ❑ Gradient Boosted Model
 - ❑ Neural Networks
 - ❑ XGBoost
- ❑ **Lesson Learned**



Lesson learned

- ❑ Prioritize team's overall strategy.
- ❑ Imputation turns out to be necessary for us (for EDA, feature importance, NN, RF).
- ❑ Use cloud computing for all team members not their own laptops (avoid crashing and uncertainties).
- ❑ Try larger grids if possible (our grids are still too sparse).
- ❑ All team members focus on one model VS each work on their own.

Lesson learned

What we did good?

- Had clear and solid plan at the beginning
- Each team member is very dedicated to the project
- Pinpoint bottlenecks early and react very fast
- Tuned large amount of parameter combinations

What we can improve?

Team shifted focus twice

-Monday - Wednesday (EDA, Imputation, Importance analysis, Brainstorming)

Thursday (Gave up imputation, tuned on complete training dataset)



Thanks!

